

# **ALLSTAR: New Challenge Problems for Static Analysis**

**Jailbreak Security Summit**

**October 11<sup>th</sup>, 2019**

**evm**

**@evm\_sec**

# A Future Vision for Software Reverse Engineering

**Imagine if Software RE  
could be faster, more  
effective and accessible**



- You have a tool that produces meaningful, descriptive labels in your disassembly of a completely new binary
- You automatically get a well-labeled software architecture diagram with a description of each module when analyzing a new binary
- People with little or no training can be given an accurate description of a new binary produced by an automated system

# Talk Outline

- What problems in software reverse engineering should we consider “solved” ?
- What are the gaps between our solved problems and our existing process?
- What problems should we work on next?
- How do we get started?

# My Background – The Embedded Systems RE World

- One “device” – but many boards/processors/firmwares
- Bare metal and RTOS environments
  - Large fully-linked binaries
  - No distinction between OS/libraries/application code
  - No clearly labeled system calls for clues
  - Vulnerability analysis: Takes much longer to get to the entry point
- Dynamic analysis is rarely possible
  - Unreliable JTAG/debugger access
  - Debug fuses
  - Debug BGA balls not pulled out on the motherboard
  - Integrated cores with no physical connection (debugged in development on dev boards)

**What problems in software  
reverse engineering should  
we consider “solved” ?**

# Solved Problems?

- Decompilation
- Function-to-function Matching
- *Combined Static & Dynamic Analysis Approaches\**

*\*since dynamic analysis is rarely possible in my world I am not qualified to evaluate the state of the toolsets (which mostly exist in other environments)*

# Decompilation

- IDA Pro,<sup>[1]</sup> GHIDRA,<sup>[2]</sup> RetDec,<sup>[3]</sup> JEB<sup>[4]</sup>
- Follow process first described by C. Cifuentes in 1994<sup>[5]</sup>
  - Syntax Analysis
  - Semantic Analysis
  - Intermediate Code Generation
    - SLEIGH/p-code (GHIDRA) <sup>[6]</sup>
    - Microcode (IDA) <sup>[7]</sup>
  - Control Flow Graph Generation
  - Data Flow Analysis
  - Control Flow Analysis
  - Code Generation
- Decompilers always produce blank code!
- Incremental improvements in decompilation will not lead to fundamental changes in the speed & effectiveness of RE


# Function Matching

- BinDiff<sup>[8]</sup> and Diaphora<sup>[9]</sup>
- Algorithm
  - Lossy compression function at the basic block level
    - Small primes product <sup>[10]</sup>
  - Graph-based comparison function for function flow graph
    - MD Index <sup>[11][12]</sup>
    - KOKA hash <sup>[13]</sup>
  - Graph-based comparison functions for call graph
- Works fairly well in most situations where library code is pre-identified
- Incremental improvements to function matching will likewise not lead to fundamental changes in the speed & effectiveness of RE

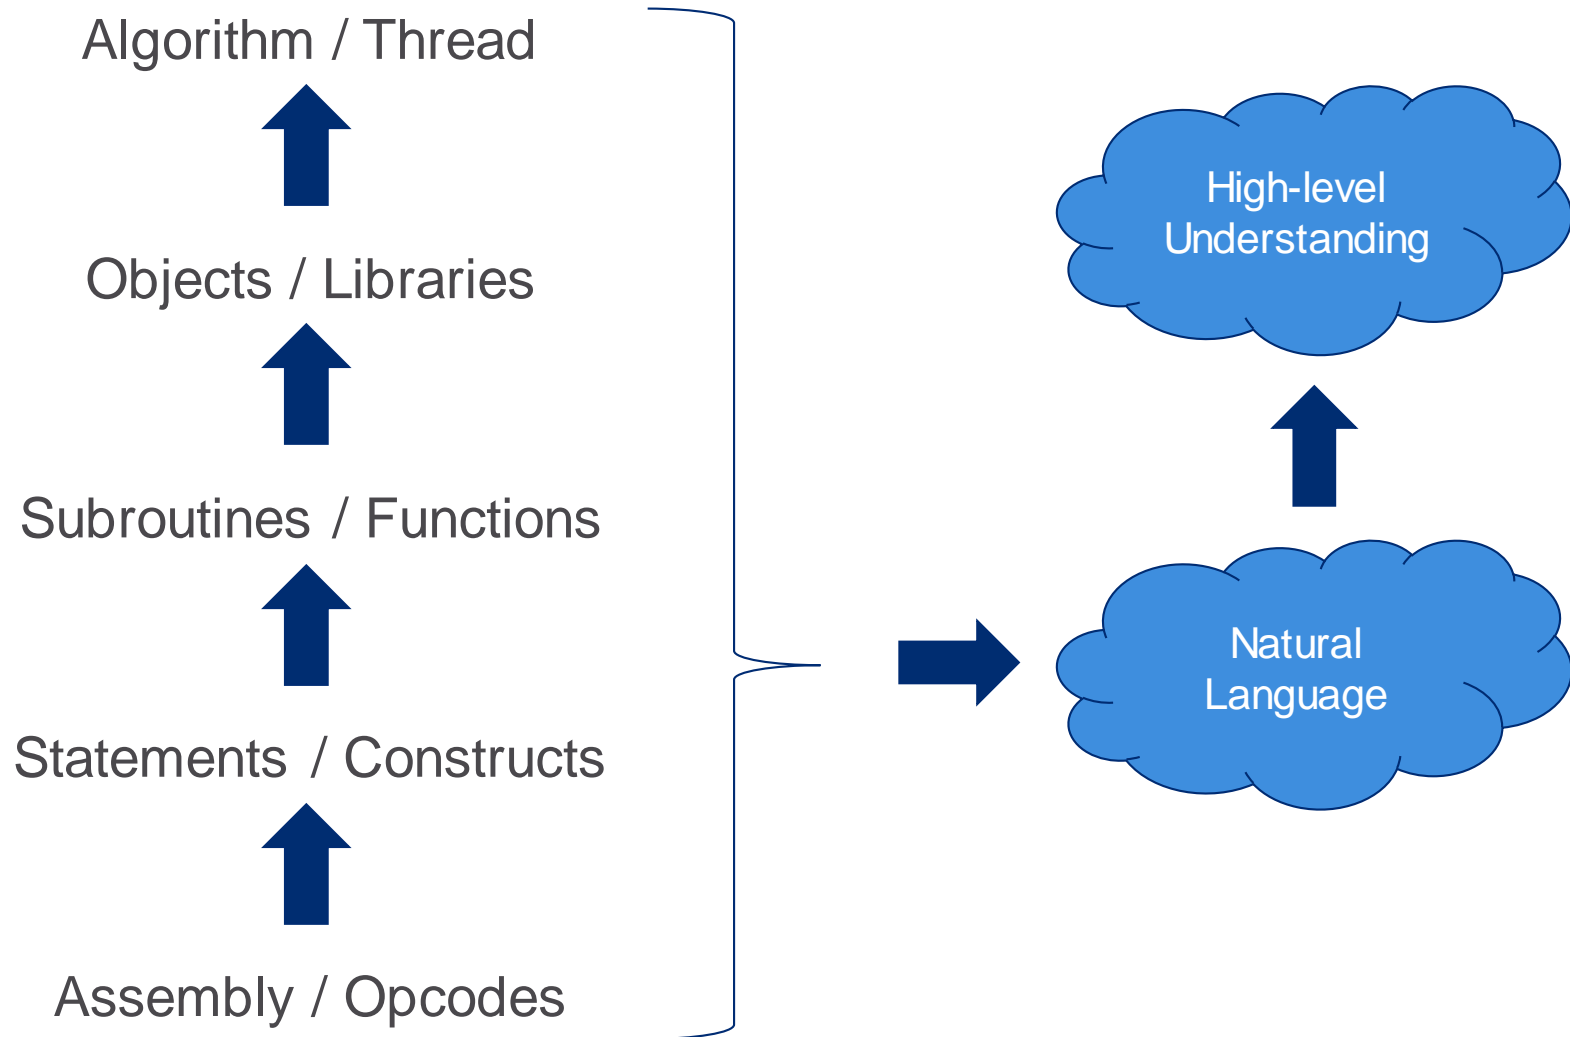


**What are the gaps between  
our solved problems and  
our existing process?**


# What is Software Reverse Engineering?



# Software Reverse Engineering Process



# Software Reverse Engineering Process



**What kind of problems  
should we work on next?**

# What's Next? – Challenge Problems

- **Variable Name Prediction** – Given a blank decompiled function, output meaningful names for the variables in that function
  - Initial research from CMU, using statistical machine translation (SMT) approach<sup>[14]</sup>
  - Inspired by work recovering identifiers in obfuscated JavaScript<sup>[15]</sup>
- **Statement Commenting** – Given a blank decompiled code statement or fragment, output comments in natural language describing it
  - What is this basic block doing?
  - Existing work in code snippet tagging/labeling<sup>[16]</sup>
  - Feed in context information from function
  - Build on variable name prediction
- **Function Summarization** – Given a blank decompiled function, output comments in natural language summarizing the function
  - Language summarization is a difficult problem due to lack of datasets
  - Existing research in source code summarization<sup>[17][18][19]</sup>
  - Needs to be built on variable name prediction, possibly statement commenting?
  - Rely on man page / documentation descriptions?

# What's Next? – Challenge Problems

- **Library/Object Organization (“CodeCut”)** – Given a fully linked binary with no symbols, locate the boundaries between the original object files – or the boundaries between related sections of code.
  - Preliminary solution using a function affinity metric and a weighted cut graph approach<sup>[20]</sup>
  - Source code available for IDA Pro<sup>[21]</sup>
- **Library Summarization** – Given an identified set of related decompiled functions, output a name or labels describing it
  - Extension of function-level summarization work
  - Extension of code snippet tagging work
- **Foundational NLP Work**
  - asm2vec<sup>[22]</sup>
  - code2vec<sup>[23]</sup>

# How do we get started?



# Assembled Labeled Library for Static Analysis Research (ALLSTAR)

- NLP research will require a large (publicly available) database of:  
Open Source Code → Compile Artifacts → Binaries (with Symbols)
- There is a TON of open source code
- There is a TON of firmware out there
- But....there are ZERO debugging symbols  
(ok almost zero)
- And \*ideally\* it would be cross-architecture...(so we avoid overtraining on x86)

# ALLSTAR Build Process

- Uses Debian “jessie” distribution
- Try to remove packages that are documentation, debug, data-only, Python, Java, etc. – leaves over 32,000 packages
- Build for all Debian architectures (x86, amd64, ARM, PPC, MIPS, s390x) using Dockcross (Docker containers with cross compilers)
- Fairly simple technical details:
  - CFLAGS: -g -fdump-final-insns -fdump-tree-gimple
  - CXXFLAGS: -g -fdump-final-insns -fdump-tree-gimple -fdump-class-hierarchy
  - DEB\_BUILD\_OPTIONS: nostrip debug
  - Build with ‘debuild’ setting architecture
- Takes 5 to 6 weeks (?) to build, 35-55 packages per hour

# ALLSTAR Stats\*

- 1.1 TB total storage
- Binaries:



- i386, MIPS and PPC numbers reflect Debian packages that properly conform to cross-arch build spec
- ARM and s390x lower due to Dockcross' cross compiler (uses a sysroot)
- ~20K cross-platform binaries is still a nice dataset

\*projected, build currently in process

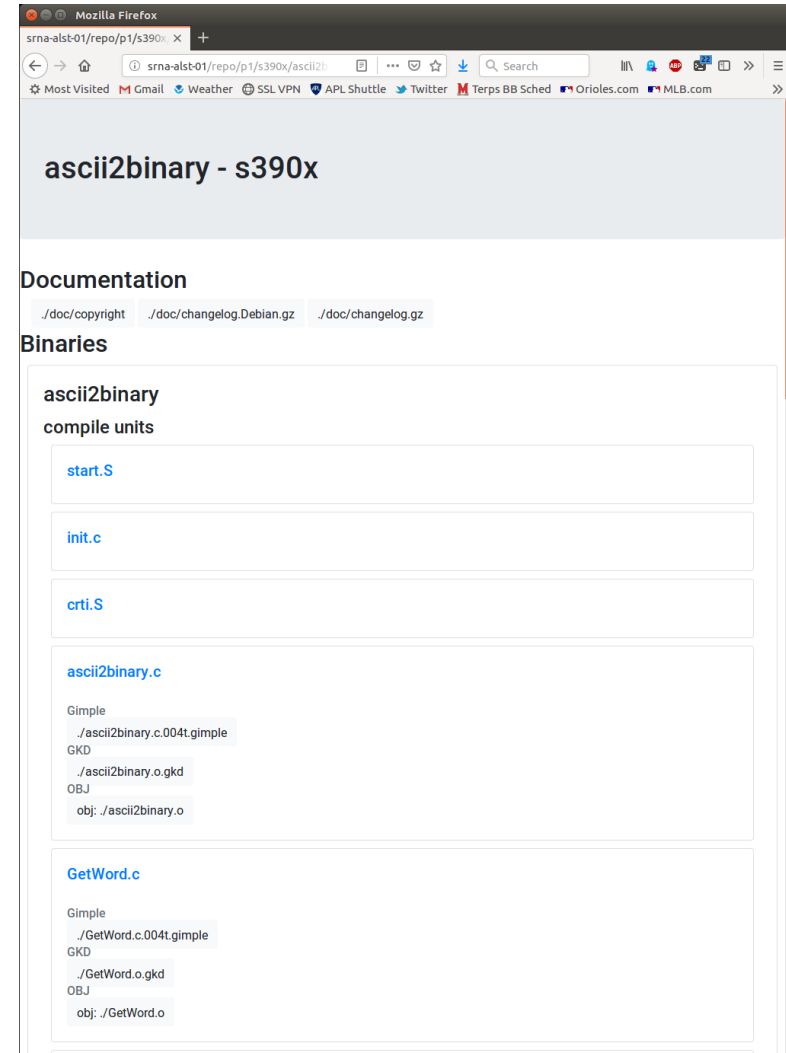
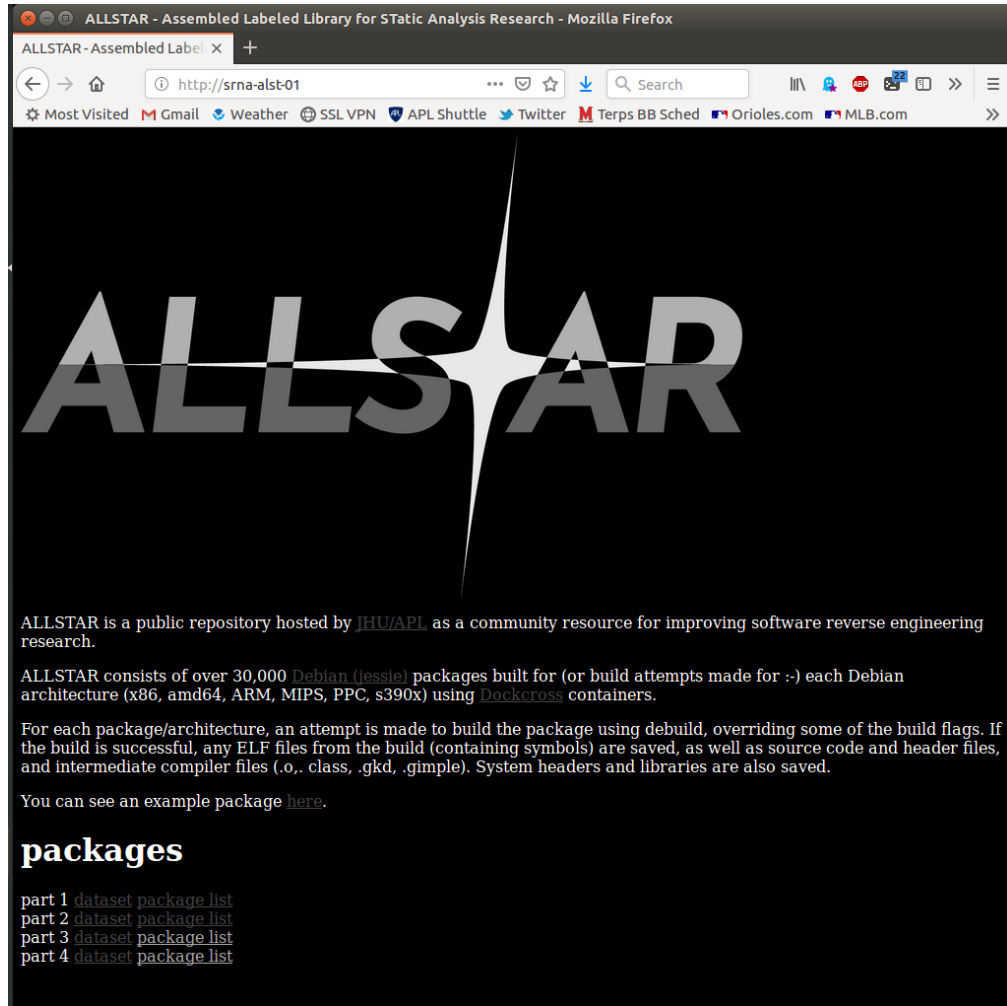
# Debian vs. Github

~32K packages (mostly C/C++)	~400K packages in C <sup>[14]</sup>
Structured build process (debbuild)	Build by trying “configure” and “make”
Packages build cross-architecture (theoretically, if they follow spec)	Less likely to build for non-amd64
Less duplication	More duplication (up to 70%! <sup>[23]</sup> )
GPL or similar open licensing	Licensing unclear / non-obvious
More serious/polished code?	A bit more random code?

# What's in ALLSTAR? A Single Data Record

- Binary (fully linked ELF, either runnable or .so)
  - Source file
    - Header files
    - GIMPLE file (like an abstract syntax tree)
    - .class file (if C++)
    - gkd file (final internal representation in RTL)
    - Object file
  - Man page
  - Symbols (in binary) – can be used to produce a .map file
  - Any system library dependencies
- Package Documentation
- HTML index for browsing
- JSON index for code parsing

# ALLSTAR – Browser Interface



# Building on ALLSTAR

- Open source / publish by end of 2019
- 2 internal research projects for FY20 planned
- Hope to see a lot of research building on ALLSTAR in the future



JOHNS HOPKINS  
APPLIED PHYSICS LABORATORY






# Special Thanks To:

- halvarflake

# Additional Thanks To:

- Igor Skochinsky
- Joxean Koret
- Joan Calvet

# Questions?



# Questions?

evm

@evm\_sec

evm.ftw@gmail.com

# References (1/4)

- 1) Hex-Rays. "IDA: About." <https://www.hex-rays.com/products/ida>.
- 2) National Security Agency. "Ghidra." <https://ghidra-sre.org/>.
- 3) Avast Software. "RetDec::Home." <https://retdec.com/>.
- 4) PNF Software. "JEB Decompiler by PNF Software." <https://www.pnfsoftware.com>.
- 5) C. Cifuentes. "Reverse Compilation Techniques." Phd diss. Queensland University of Technology, 1994.
- 6) National Security Agency. "Ghidra Pcode." *Ghidra source code*. <https://github.com/NationalSecurityAgency/ghidra/tree/master/Ghidra/Framework/SoftwareModeling/src/main/java/ghidra/pcode>.
- 7) I. Guilfanov. "Decompiler internals: microcode." *RECON 2018*. Brussels, Belgium. <http://www.hex-rays.com/products/ida/support/ppt/recon2018.ppt>.

# References (2/4)

- 8) zynamics. “BinDiff.” <https://www.zynamics.com/bindiff>.
- 9) J. Koret. “Diaphora.” <https://diaphora.re>.
- 10) T. Dullien, and R. Rolles. "Graph-based comparison of executable objects." *SSTIC 5, no. 1*. 2005.
- 11) T. Dullien, et al. “Automated Attacker Correlation for Malicious Code.” *Information Systems and Technology Panel (IST) Symposium*. Talinn, Estonia. 2010.
- 12) H. Flake and S. Porst. “ShaREing Is Caring.” *CANSEC West 2010*. Vancouver, Canada. 2010.  
<https://www.slideshare.net/cblichmann/shareing-is-caring-3856188>.
- 13) C. Karamitas, and A. Kehagias. "Efficient features for function matching between binary executables." *In 2018 IEEE 25th International Conference on Software Analysis, Evolution and Reengineering (SANER)*, pp. 335-345. IEEE, 2018.

# References (3/4)

- 14) A. Jaffe, et al. “Meaningful variable names for decompiled code: A machine translation approach.” In *Proceedings of the 26<sup>th</sup> Conference on Program Comprehension* (pp.20-30). ACM.
- 15) B. Vasilescu, et al. “Recovering clear, natural identifiers from obfuscated JavaScript names,” in *Joint Meeting of the European Software Engineering Conference and the Symposium on the Foundations of Software Engineering*. 2017, pp.683-693.
- 16) B. Gelman et al. “A language-agnostic model for semantic source code labeling.” In *Proceedings of the 1<sup>st</sup> International Workshop on Machine Learning and Software Engineering Symbiosis*, pp.36-44. ACM, 2018.
- 17) J. Moore et al. “A Convolutional Neural Network for Language-Agnostic Source Code Summarization.” *arXiv preprint arXiv:1904.00805*, 2019.
- 18) S. Iyer et al. “Summarizing source code using a neural attention model.” In *Proceedings of the 54<sup>th</sup> Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, 2016.

# References (4/4)

- 19) X. Hu et. Al. “Deep code comment generation.” *In Proceedings of the 26<sup>th</sup> Conference on Program Comprehension*, pp.200-210. ACM, 2018.
- 20) evm. “A Code Pirate’s Cutlass: Recovering Software Architecture from Embedded Binaries.” *Shmoocon 2019*. Washington, DC. 2019.
- 21) evm. “CodeCut: Locating Object File Boundaries in IDA Pro.” <https://github.com/JHUAPL/CodeCut>.
- 22) S. D’Antoine. “asm2vec: Binary Learning for Vulnerability Discovery.” *Jailbreak Security Summit 2018*. Laurel, MD. 2018.
- 23) U. Alon et al. “code2vec: Learning distributed representations of code.” *Proceedings of the ACM on Programming Languages* 3, 2019.
- 24) P. Martins, et al. “50K-C: A dataset of compilable and compiled, Java projects.” In *2018 IEEE/ACM 15<sup>th</sup> International Conference on Mining Software Repositories (MSR)*, pp 1-5. IEEE, 2018.